

AMAZON AI CYBER SECURITY HACKATHON

Round 1 Participant Brief

Build the pitch deck and architecture for a final product that makes India’s digital financial ecosystem safer from AI-enabled threats.

ROUND 1 OUTCOME Top 30 teams move to the in-person 48-hour finale.	YOUR SUBMISSION One PDF: a pitch deck plus an architectural overview.	WHAT MATTERS A clear, safe, technically credible product that you can build.
--	---	--

Choose one challenge statement

01	Voice Cloning Fraud Detection	Design a safe, real-time way to identify voice-cloning fraud in banking or payment workflows while protecting privacy and limiting disruption to legitimate users.
02	AI-Driven Scam Pattern Recognition	Detect emerging AI-enabled scam workflows in UPI, banking, or wallet contexts and deliver timely, useful alerts without creating alert fatigue.
03	Deepfake Detection for Financial Communications	Help people and institutions verify financial communications, detect manipulation, and create an actionable path for review, reporting, or response.

Select exactly one track. You may mention future extensions to other tracks, but your Round 1 submission will be evaluated against your selected track.

Your Submission

Submit one PDF containing both items. A reviewer must be able to understand your idea without a live presentation or an external link.

1. Pitch deck

Maximum 10 slides, excluding an optional appendix.

- Team, selected challenge statement, and one-sentence solution
- User, threat scenario, and the harm you are solving
- Current journey, product flow, and solution differentiation
- Core detection or decision-making approach
- Data, evidence, validation plan, and key assumptions
- Privacy, security, safety, and misuse safeguards
- Deployment pathway and 48-hour prototype plan

2. Architectural overview

Maximum 3 pages or 3 slides.

- One legible end-to-end system diagram
- Inputs, data sources, consent, and provenance
- Components, models or rules, storage, interfaces, and outputs
- Detection or decision flow, alerts, and escalation
- Human review, notifications, reporting, and evidence path
- Security, privacy, and adversarial-risk mitigations
- Performance assumptions and validation approach

Submission rules

- Teams may have 2 to 4 student members and submit one entry.
- Use only synthetic, public, consented, or properly authorised data. Do not include confidential bank, customer, or platform data.
- Disclose AI-generated content, datasets, models, and third-party assets in an appendix or slide note.
- Incomplete submissions may be ineligible. Reviewers will score the submitted PDF only.

What Reviewers Will Evaluate

The rubric rewards substance over slide polish. Your score is out of 100 points.

Criterion	Weight	What strong submissions show
Problem fit and user understanding	20	A specific, India-relevant threat scenario; clear primary user; credible stakeholder journey; a sharply bounded problem.
Technical architecture and detection logic	20	A legible end-to-end design, defensible detection or verification logic, defined inputs and outputs, and realistic boundaries.
Solution value and differentiation	15	A clear before-and-after workflow, meaningful user value, and more than a generic AI detection claim.
Feasibility and finale execution plan	15	A narrow, demonstrable MVP; named assumptions and dependencies; and a way to prove progress in 48 hours.
Safety, privacy, and responsible AI	15	Concrete safeguards, a safe escalation path, and awareness of false positives, bias, misuse, and adversarial adaptation.
Adoption, integration, and scale	10	A practical integration pathway, operating owner, rollout sequence, and attention to latency or operational friction.
Communication quality	5	A concise, well-structured, legible PDF with a clear narrative and well-labelled architecture.

Build for Trust

Financial-safety products can create harm if they are wrong, intrusive, or easy to abuse. Make your safeguards visible in both the deck and architecture.

Design question	Show this in your submission
What data do you need?	Name the inputs, why each is necessary, how consent or authorisation works, and what you do not collect.
What happens when the system is wrong?	Explain false-positive and false-negative handling. Avoid automatic blocking or accusations based on a single weak signal.
Who can act on the result?	Define the user, the action they can take, and any human review or escalation needed before high-impact action.
How will users understand the result?	Use clear, actionable messages. A detector score alone is not an explanation or proof.
How will the product resist misuse?	Consider fraudster adaptation, adversarial inputs, privacy abuse, and unauthorised access to sensitive evidence.

Track-specific technical cues

- **Voice cloning:** Consider audio, behavioural, or call-context signals, customer friction, consent, and a safe evidence path.
- **Scam patterns:** Define the scam workflow, available metadata or behavioural signals, timely action, alert fatigue, and transaction privacy.
- **Deepfakes:** Consider provenance and forensic signals, explanation for non-technical reviewers, evolving manipulation methods, and a credible verification or reporting path.

Final Check Before You Submit

Use this checklist. A reviewer should be able to find each answer quickly.

01	The problem	Can someone identify the primary user, harmful moment, and exact threat pattern in under a minute?
02	The product	Does the deck show what happens from input to alert, decision, verification, or evidence record?
03	The architecture	Is there a legible system diagram with clear components, data boundaries, and outputs?
04	The proof plan	Have you stated your assumptions, how you will validate them, and what you can demonstrate in 48 hours?
05	The safeguards	Have you addressed privacy, false positives, bias, user communication, human review, and misuse risks?

Submit one clear PDF.

Pitch deck: maximum 10 slides. Architectural overview: maximum 3 pages or 3 slides. Select one challenge statement and design a product that is useful, buildable, and safe.

Questions about the submission format should be directed to the organising team before the application deadline.